

Multinomial Mixture Model (Multiple trials)

Setting:

- $K = 3$ possible clusters (classes) (π denotes the mixing proportion)
- For each k , there are three possible outcomes $J = 6$ (θ_k denotes multinomial parameters for cluster k and θ_{kj} denotes the probability of having outcome j in cluster k). We have 10 trials per sampling.
- Use Dirichlet distribution prior for π and θ_k with non-informative prior

Summary on data generating process

$$\pi_i \sim \text{Dirichlet}(1)$$

$$\theta_k \sim \text{Dirichlet}(1)$$

$$x_i | z_i = k, \theta_k \sim \text{Mult}(10, \theta_k)$$

```
# Define mixing proportion
set.seed(0)
pi_true = c(0.3, 0.2, 0.5)
theta_1_true = c(0.33, 0.33, 0.34, 0, 0, 0)
theta_2_true = c(0, 0, 0.1, 0.6, 0.2, 0.1)
theta_3_true = c(0.1, 0.1, 0, 0, 0.1, 0.7)

theta_true = rbind(theta_1_true, theta_2_true, theta_3_true)

# Create simulated data
n = 1000
class_i = rmultinom(n, size = 1, prob = pi_true)
x = c()
for (i in 1:n) {
  x = cbind(x, rmultinom(1, size = 10, prob = theta_true[class_i[, i]==1,]))
}

# Prior Parameter
cluster_num = 3
category_num = 6
a_pi = rep(1, cluster_num) # For Dirichlet distribution for pi
a_theta = rep(1, category_num) # For Dirichlet distribution for theta
```

Blocked Gibbs Sampling here consists of three major steps (corresponding to 3 parameters of interest z_i, θ, π)

1. Sample cluster indicator z_i for each x_i from full conditional multinomial distribution

For each i :

$$P(z_i = k | x_i, \theta, \pi) = \frac{\pi_k P(x_i | z_i = k)}{\sum_k \pi_k P(x_i | z_i = k)}$$

2. Sample θ from the updated (posterior) Dirichlet distribution for each of the 4 clusters

For each k :

$$\theta_k | x, z \sim \text{Dirichlet}(1 + n_{k1}, 1 + n_{k2}, 1 + n_{k3})$$

where n_{ki} is the number of observations found in cluster k that is of category i

3. Sample π from the updated (posterior) Dirichlet distribution to obtain the new mixing proportion

$\pi \mid x, z \sim \text{Dirichlet}(1 + n_1, 1 + n_2, 1 + n_3, 1 + n_4)$

where n_k is the number of observations found in cluster k

```
# Blocked Gibbs Sampling
set.seed(1)
# Initialize parameters
pi = c(0.33, 0.33, 0.34)
theta = c()
for (i in 1:cluster_num) {
  theta = rbind(theta, rdirichlet(1, a_theta))
}

# Initialize the sampling matrix
sample_pi = c()
sample_pmf = c()
for (round in 1:2000) {
  # Step 1: Sampling cluster indicator
  z = c()
  for (i in 1:dim(x)[2]) {
    # Calculate the full conditional probability of belonging to cluster k
    fullcon_zi = pi*rowProds(t(t(theta)^x[,i]))/
      sum(pi*rowProds(t(t(theta)^x[,i])))
    z = cbind(z, rmultinom(1,1,fullcon_zi))
  }
  # Step 2: Update theta
  for (k in 1:length(pi)) {
    if (is.null(dim(x[, z[k,]==1]))) {
      # only one member of no member
      if (length(x[, z[k,]==1] == 0)) {
        # No member
        nk = rep(0, category_num)
      }else{
        # One member
        nk = x[, z[k,]==1]
      }
    }else{
      # More than one member
      nk = apply(x[, z[k,]==1], MARGIN = 1, FUN = sum)
    }
    # Sample theta using full conditional distribution
    theta[k,] = rdirichlet(1, a_theta + nk)
  }
  # Step 3: Update pi
  n = apply(z, MARGIN = 1, sum)
  pi = rdirichlet(1, a_pi + n)
  sample_pi = rbind(sample_pi, pi)

  # Record pmf
  sample_pmf = rbind(sample_pmf, apply(as.vector(pi)*theta, MARGIN = 2, sum))
}

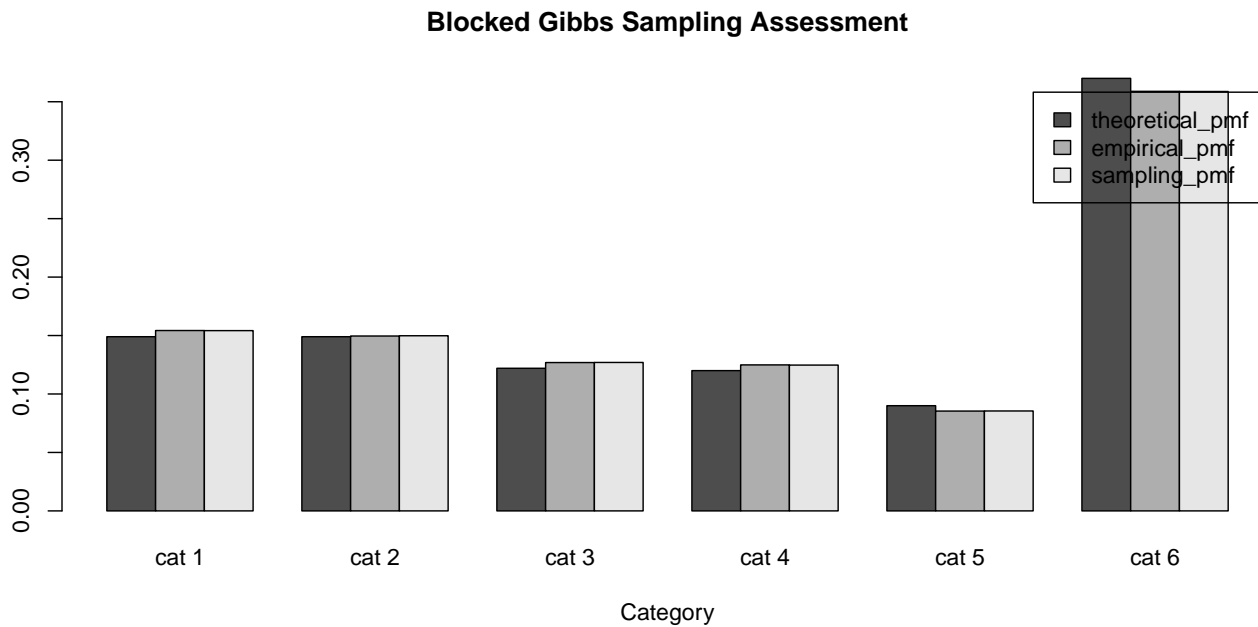
# Model checking with empirical data
theoretical_pmf = apply(as.vector(pi_true)*theta_true,MARGIN = 2, sum)
```

```

empirical_pmf = apply(x, MARGIN = 1, mean)/10
sampling_pmf = apply(sample_pmf, MARGIN = 2, mean)

df = rbind(theoretical_pmf, empirical_pmf, sampling_pmf)
colnames(df)<- c('cat 1', 'cat 2', 'cat 3', 'cat 4', 'cat 5', 'cat 6')
barplot(df, xlab = 'Category',
        beside = TRUE, legend = TRUE, main = 'Blocked Gibbs Sampling Assessment')

```



```

# Convergence checking
matplot(1:dim(sample_pmf)[1], sample_pmf, type = 'l',
        ylab = 'Prob', xlab = 'trials',
        main = 'Checking stability of marginal pmf', ylim = c(0,0.4))

```

```

matplot(1:dim(sample_pi)[1], sample_pi, type = 'l',
        ylab = 'Prob', xlab = 'trials',
        main = 'Label Switching?', ylim = c(0,1))

```


```

apply(sample_pi, MARGIN = 2, mean)

```

```
## [1] 0.3110584 0.4870247 0.2019169
```

Checking stability of marginal pmf



Label Switching?

